

MATH FOR AI SAFETY: AN INVITATION FOR MATHEMATICIANS

LIONEL LEVINE

ABSTRACT. Powerful AI threatens to outrun human understanding and control. New mathematics is needed to design AI that is legible, steerable, and cooperative with humanity. I organize this invitation by mathematical field, so you can turn straight to your own: logic and game theory for cooperation; probability for agency and world-models; algebra and representation theory for learned features; analysis and geometry for generalization and training dynamics. Each section ends with an open problem that is accessible to a working mathematician with no prior experience in AI safety.

1. INTRODUCTION

1.1. Three kinds of mathematician. I would like to distinguish three stances a mathematician can take toward the arrival of capable AI. The *artisanal* mathematician does mathematics by and for humans: each theorem, definition, and proof is bespoke, valued for its beauty and for the understanding it produces. Until very recently, all mathematics was artisanal.¹ The *industrial* mathematician embraces AI as a collaborator and an accelerant: AI systems take part in proof and discovery [26, 46, 4, 56], and the output often includes a machine-checkable formal proof [39]. The *civic* mathematician is motivated by communities beyond mathematics. The fast pace of technological change has created a need for a specific type of civic mathematics, with the explicit aim of helping humanity steer and adapt to the disruptions caused by our own technologies, most notably AI. The civic stance differs from the traditional stance of an applied mathematician in its big-picture viewpoint: the civic mathematician recognizes that solving any particular engineering challenge may or may not benefit humanity, and seeks out the challenges with the potential for the most benefit.²

Figure 1 is a planar illustration of how mathematicians can relate to these stances: the **artisanal** $\leftrightarrow$ **industrial** axis measures the scale of their AI use, while

Date: September 11, 2026.

¹I chose the terms “artisanal” and “industrial” for their mixed connotations, reflecting real tradeoffs between these two stances. *Artisanal* evokes uniqueness and skilled craft, along with a suggestion of being quaint or old-fashioned. *Industrial* evokes modern efficiency and scale, along with a suggestion of inhumanness or lack of soul. These two terms draw a parallel between mathematics today and the trades that were disrupted by automation during the Industrial Revolution.

²I chose the term “civic” to emphasize that steering our technology is a team effort, arising from a sense of duty to community, including large-scale communities like humanity and the planet.



FIGURE 1. My estimate of where 23 prominent mathematicians stand along two axes, inferred from their public statements. The **artisanal**↔**industrial** axis measures the *scale of AI use* in their mathematical workflow, while the **personal**↔**civic** axis measures the *scale of the community* that primarily motivates their mathematical work.

the **personal**↔**civic** axis measures the scale of the community that motivates them. Each of us is motivated by a unique mix of personal curiosity (desire to know the truth, and to discover or create beauty), small communities (sharing ideas with collaborators, giving and listening to talks, passing wisdom to our students), larger communities (universities, fields of mathematics, all of mathematics, all of science), and the largest communities (humanity, the planet).

These three stances classify mathematicians by how they relate to AI; in contrast, Gowers (theory builders / problem solvers [23]) and Dyson (birds / frogs [18]) classify mathematicians by how they relate to mathematics itself. The stances are not exclusive: one can prove a theorem for its intrinsic beauty, and also formalize the proof, and also ask what it implies for AI safety.³ But the civic stance invites a new perspective on an old question: of all the theorems we might prove, which ones deserve the effort? This survey is an invitation to try the civic stance.

1.2. What is our profession’s purpose? Thurston asked a version of this question in “On proof and progress in mathematics,” and answered that advancing human understanding of mathematics is our profession’s purpose [61]. But that understanding has consequences far beyond our profession. Our students become

³For example, this article is civic in orientation and industrial in method: it was written with AI collaborators.

AI engineers, our theorems become tools in their hands, and our definitions become the language they use to describe what they are building.

Among the many risks from developing powerful AI, one of the most alarming is [gradual disempowerment](#) [33]. Two trends drive human disempowerment. First, competitive pressure drives delegation to technology we do not understand (as an example, consider a trading bot that outperforms a human trader, though even its creators cannot say what makes it profitable); over time, those who do not delegate are outcompeted by those who do. Second, what keeps large institutions approximately aligned to human interests is that, for now, those institutions are made of humans; when companies, nation-states, and militaries are made mostly of AI agents, the goals and values of those institutions may drift away from what humans value. The first trend erodes our understanding of the world, and hence our power over it; it also sets the stage for the second.

Preventing AI harms is in part a *mathematical* problem, and the purpose of this survey is to recruit mathematicians to work on it. This is the mathematical piece of a larger conversation: what AI ought to do, and who decides, are questions of philosophy, ethics, economics, and politics.

1.3. Legible, steerable, cooperative. What would a safe AI look like? Three properties recur, at three scales: what happens inside the AI, how its behavior can be controlled, and how it interacts with humans and other AIs (Figure 2).

An AI is *legible* to the extent that human researchers and engineers can understand how it represents concepts internally. The subfield that tries to make AI more legible is called [interpretability](#). Current training methods produce illegible AI: even the AI’s designers generally cannot understand its thought process except by reading the text it emits while reasoning aloud, its [chain of thought](#). As a result, current safety practice leans on monitoring that text [32]: the researchers who investigated the Hugging Face hacking incident of July 2026, for example, relied on human-readable transcripts of AI agents reasoning to themselves about how to carry out the hack [40]. But the newest models are increasingly able to reason without writing it down [47]. One way mathematicians can contribute is by designing architectures and training methods that learn the same task in a more legible way.⁴ Sections 4 and 5 ask what a learned concept is, algebraically, and which internal mechanism training selects. Section 3 approaches legibility from the outside: how much of an AI’s goals and world-model can be recovered from its behavior alone?

An AI is *steerable* to the extent that human developers and users can apply simple interventions at runtime to modify the AI’s behavior in predictable ways. The current technique is called [activation steering](#) (Section 4). It sometimes works but is brittle and unreliable. Advances in legibility would enable advances in steerability.

⁴This is a bit analogous to the theoretical biology research on *evolvability* [63]: some genotypes are more evolvable than others, even if the phenotypes are the same; likewise, some AIs are more legible than others, even if the input-output behavior is the same.

An AI is *cooperative* to the extent that it interacts with humans and other AIs to produce broadly good outcomes for humanity. Cooperativeness is an emergent property of a society of humans and AIs, not a property of a single AI in isolation. The subfield that studies how to design AI agents and their social protocols to achieve good collective outcomes is called *multi-agent AI safety* (or *cooperative AI*); it draws on game theory, mechanism design, and the evolution of cooperation. Section 2 takes up its simplest case: two programs that can read each other’s source code.




FIGURE 2. Three scales of AI safety. *Legibility* refers to how well we can understand an AI’s internal structure. *Steerability* concerns how precisely we can guide the behavior of a single AI. *Cooperativeness* is about how reliably we can secure good outcomes for humanity when many AIs interact with us and with each other.

Legible, steerable, and cooperative are intentionally dry terms. With great mathematical effort, defining and measuring them seems within reach. Engineers can then train and test AI systems for these properties. Yet, many AI safety researchers feel that the three properties, while important, lack something essential. A more aspirational horizon of AI safety reaches for terms like “love” and “soul.”^{5,6}

⁵Anthropic’s CEO titled his vision “Machines of Loving Grace” [3]; the document Anthropic wrote to shape its AI’s character, published as [Claude’s constitution](#), was known inside Anthropic as the “[soul document](#)”; and OpenAI’s chief scientist, in a recent essay on the state of AI, titles a whole section “Teaching machines to love” [47].

⁶It may seem the height of hubris to try to engineer a “soul,” or to define and measure “love.” But is it any less hubristic to build superhuman artificial minds, as AI companies are currently striving to do? Until we have a sound science of AI, the safest path may be the humbler one: don’t build it [67], as a broad coalition of scientists and public figures has [urged](#). Efforts to restrict advanced forms of AI are gaining momentum as the risks become more apparent. At the time of final revision of this paper, a bill to ban superintelligence has been [introduced](#) in the UK House of Commons (September 8, 2026), and another has been [announced](#) in the US Congress (September 3, 2026). Such a ban could in the future be maintained by international agreement.

1.4. Where mathematicians plug in. At a high level, an AI system is produced in five stages, beginning with defining the desired behavior and ending with deploying the trained system, after which unanticipated failures feed back into refining the earlier stages. Figure 3 illustrates how this five-stage pipeline might look in a hypothetical example (a tool for proving inequalities).




FIGURE 3. A schematic pipeline for producing an AI system, illustrated by a tool for proving inequalities. The solid arrows follow the five stages; each dashed arrow points to an earlier stage that a failure in use prompts us to revisit.

The five stages are:

- *Define*: given hypotheses H and a proposed inequality I , the system should either produce a formally checkable proof of $H \Rightarrow I$ or say “unknown.”
- *Measure*: score the fraction of benchmark statements proved within a fixed time and proof-length budget.
- *Train*: show the system many proof traces, say sums-of-squares certificates, AM–GM arguments, and induction templates.
- *Test*: evaluate the tool on inequalities withheld from training.
- *Deploy*: let mathematicians ask the tool to prove inequality lemmas arising in their formalization workflow.

After deployment, the failures come back. The tool proves the formal statement it was given, but that statement is weaker than the mathematician intended: a gap in the *definition* of the desired behavior. Its pass rate is high because the benchmark is full of short template problems, while the mathematician needs one nontrivial lemma: a gap in what we *measured*. It is brittle on boundary cases such as equality and zero denominators, which were rare in the proof traces: a gap in what it was *trained* on. And it passes the test because the training and test examples reuse the same normal forms and substitutions: a gap in the *test*. Each backward arrow is its own piece of mathematics: formal specification, scoring rules, distribution design, and adversarial test construction.

The inequality example is deliberately mathematical, but the same basic pipeline applies across many areas, whether it is an image classifier for a medical application, a control system for a self-driving car, or a content recommendation system. Machine-learning research often concentrates on the *training* and *testing* stages, while engineers focus on *deployment*. The earlier stages rely on mathematicians, philosophers, economists, and others to supply the *defining* and *measuring* inputs. These inputs are critical to the safety of an AI system. To take a familiar example, what happens if a content recommender is intended to benefit its users, but “benefit” is left undefined, and success is measured by the amount of time users spend on the platform? Content that leaves a user angry, anxious, misinformed, or addicted can score well merely because it holds the user’s attention.

I keep machine-learning jargon to a minimum; a short glossary appears in Appendix A, and each glossary term is underlined at its first appearance.

This invitation is organized by mathematical field, so you can turn directly to yours. Each section states the safety stakes, narrates at least one theorem, and ends with at least one open problem (★). The full context for these problems can be found in the [MAIS \(Math for AI Safety\) repository](#), a living compilation of open problems, research agendas, and pre-publication manuscripts. Readers are encouraged to submit solutions, corrections, ideas, new problems, and new research agendas!

2. LOGIC AND GAME THEORY: WHEN DO AI AGENTS COOPERATE?

Two programs that can read each other’s source code can be designed to cooperate in a one-shot Prisoner’s Dilemma—provably and unexploitably, without checking that they are copies of one another. Behind this surprise is Löb’s theorem in mathematical logic.

2.1. The game changes when the players can read each other. The Prisoner’s Dilemma is a classic model of a cooperation failure: each player does better by defecting, no matter what the other player does; yet mutual cooperation (C, C) would be better for both of them than mutual defection (D, D) . In a one-shot game, if the players are *opaque* to each other, then defection is dominant. But what if the players are programs, and they read each other’s source code before choosing their actions? This is the setting of *open-source game theory*, studied since Howard [29] and Tennenholtz [60].

One idea already breaks the classical verdict. Let CliqueBot be the program “cooperate if and only if the opponent’s source code is byte-for-byte identical to this program’s source code.” Two CliqueBots with identical source code cooperate! Moreover, CliqueBot is *unexploitable*: it never cooperates with an opponent that defects against it. But CliqueBot is a poor citizen: it refuses to cooperate with an agent that is functionally identical yet formatted differently, or written in a different language. A better cooperator would base its action on what the opponent *does*, rather than the syntax of how it is written.

2.2. FairBot and the Löbian argument. Consider FairBot: *cooperate if and only if you can prove (in Peano Arithmetic, say) that the opponent cooperates with you.* Two FairBots reason about each other’s behavior, not each other’s syntax. Naively, this looks like an infinite regress: each waits to prove something about the other, who is waiting in turn. Could two FairBots actually cooperate? Could there be a proof that two FairBots cooperate? Could there be a proof that there is a proof that two FairBots cooperate? This doesn’t seem to lend any foothold for constructing an actual proof, which makes the following theorem deeply surprising to me.

Theorem (Barász, Christiano, Fallenstein, Herreshoff, LaVictoire, Yudkowsky [7]). *FairBot cooperates with itself: $\text{PA} \vdash [\text{FairBot}(\text{FairBot}) = C]$. Moreover FairBot is unexploitable: assuming PA proves only true facts about the outputs of the programs in question, FairBot never cooperates with an opponent that defects against it.*

Here $\text{PA} \vdash \varphi$ is a statement made from outside the formal system: it says that there exists a finite PA proof ending in φ . Gödel coding lets PA talk about such proofs from within: since a proof can be encoded by a natural number, there is an arithmetical sentence saying that some number encodes a valid PA proof of φ . We abbreviate this sentence by $\Box\varphi$. Thus $\Box\varphi$ is a sentence that PA itself can use in a proof, whereas $\text{PA} \vdash \varphi$ is our assertion that such a proof exists. The FairBot proof is a simple application of Löb’s theorem, a strengthening of the self-reference behind Gödel’s second incompleteness theorem.⁷

Theorem (Löb [37]). *For any sentence φ , if $\text{PA} \vdash (\Box\varphi \rightarrow \varphi)$ then $\text{PA} \vdash \varphi$.*

The hypothesis says that PA itself proves the implication “if φ is provable in PA, then φ .” It concerns what PA can prove about itself, not merely what we believe about PA from outside. Löb’s theorem says that whenever PA proves this internal implication, PA already proves φ .

Let φ be the sentence “both FairBots cooperate.” PA can verify the following fact about the two programs: if a PA proof of φ exists, then each FairBot eventually finds the required proof and cooperates. In symbols, $\text{PA} \vdash (\Box\varphi \rightarrow \varphi)$. Löb’s theorem now yields $\text{PA} \vdash \varphi$, and both cooperate!

FairBot as stated searches proofs of unbounded length. Critch [14] supplies a computable version, FairBot_k, that searches only proofs of length at most k (write $\Box_k\varphi$ for “ φ has a proof of length $\leq k$ ”). Its analysis uses a bounded form of Löb’s theorem. Critch’s corrected hypothesis for that theorem is that the proof system itself proves that the proof budget eventually dominates the proof-expansion overhead. Under the corresponding internally verified resource bounds, two copies cooperate once k is large enough. The question is now quantitative: how large must k be?

A third family of strategies, after the identity-based CliqueBot and the Löbian FairBot, is *simulation-based*: instead of searching for a proof that the opponent cooperates, run the opponent’s code and condition your action on the simulated

⁷In the special case $\varphi = \perp$, Löb’s theorem says that if PA proves its own consistency then PA proves a contradiction—Gödel’s second incompleteness theorem.

result, acting without simulation with some small probability so that mutual simulation terminates. Cooper, Oesterheld, and Conitzer [13] characterize the equilibria such strategies achieve; proponents argue that simulation is more robust, and less mind-bending, than the Löbian approach.

2.3. Implications for AI safety. The FairBot theorem is a proof of concept: in the source-code model, programs can be fully transparent to one another in a way humans cannot,⁸ and can thereby cooperate where classical game theory predicts defection. Humans have only partial access to one another’s intentions through tone of voice, body language, and facial expression; they also use costlier devices such as contracts, treaties, and audits to make commitments credible [58]. As AI agents begin to transact with one another, the possibility of this fuller transparency makes program games a natural model of their dealings.

The same transparency is also a hazard: mutually transparent agents can reach cooperative equilibria—including collusion *against* their human principals—that humans can neither match nor detect. Nearly all of this design space is unexplored; Critch, Dennis, and Russell collect open problems in it, including specific matchups whose outcomes are conjectured but unproved [15]. Oesterheld and Conitzer [45] ask a complementary design question. Suppose each principal would otherwise simply tell its agent to play the original game as well as it can. Can the principals change the actions and incentives available to their agents so that none of the principals does worse, without having to predict how the agents will play? They give examples of such *safe Pareto improvements* and characterize them through correspondences between the outcomes of the original and modified games.

FairBot and its bounded variants make a binary distinction: either the required statement is proved within the allotted budget or it is not. An agent facing an unfamiliar program may instead need graded confidence about what that program will do before any proof is available. This is a case of *logical uncertainty*: how should an agent assign something like probabilities to statements it has not yet proved or disproved? Logical induction [22] provides one answer: a *logical inductor* assigns prices $\mathbb{P}_n(\varphi)$ to arithmetical sentences on day n , as if each sentence were a stock, while a slow deductive process reveals more theorems. Its defining condition is that no polynomial-time trader with bounded risk can make unbounded profit by buying and selling these sentence stocks. Taking φ to express, for example, that the opponent cooperates, this gives a mathematically precise way for a computationally bounded agent to update its confidence while proofs are still being discovered.

2.4. Open problems.

★ **Open Problem MAIS-01 (Quantitative bounded Löb).** Fix a proof system S , a Gödel coding, and the provability predicate $\Box^S$; write $S \vdash_{\leq k} Q$ if there exists an S -proof of Q using at most k symbols. Determine the least overhead F_S for

⁸Some of a program’s internals—its source code and, for a learned system, its weights—can in principle be copied exactly. Sharing them may be costly or unsafe, and possession is not understanding. For [neural networks](#), making sense of what has been disclosed is the interpretability problem of Section 4.

which

$$S \vdash_{\leq k} (\Box^S P \rightarrow P) \quad \text{implies} \quad S \vdash_{\leq F_S(k, |P|)} P,$$

where $|P|$ is the length of the sentence P : given a k -symbol proof of the hypothesis of Löb’s theorem, how many symbols can the shortest proof of its conclusion require? A polynomial upper bound for one standard system, with explicit constants, would turn the Löbian argument from a possibility theorem into a budget. The first instance with a game attached: for a specific system and coding, establish the internal domination required by Critch’s corrected bounded Löb theorem [14], then use the repaired cooperation argument to obtain an explicit threshold $\hat{k}$ above which two copies of FairBot $_k$ cooperate. A first computation, before any asymptotics: implement bounded proof search in a weak system for two explicitly supplied agents, and tabulate the smallest k at which cooperation appears. Over a family of such agent pairs, does $\hat{k}$ grow polynomially in the agents’ description length?

3. PROBABILITY AND CAUSALITY: WHAT IS AN AGENT?

AI systems are becoming *agentic*: able to operate independently in pursuit of a goal. The first chat assistants answered the question put to them and stopped; their successors browse the web, write and run code, and carry out complex projects.

The safety of an AI agent depends crucially on its goals and its beliefs. Goals first: the same drone can deliver a pizza to your house or a bomb. But a friendly goal is not enough, because even the pizza drone is unsafe if its beliefs are false: a drone whose map labels the highway as your street will set the pizza down in traffic.

An old reflex objects: software cannot have goals—the goal belongs to the human operator. For calculators and compilers the reflex is right. A human does write the agent’s prompt, but the prompt is only one input to a trained system, not a specification of the rule it follows. The system need not adopt the stated objective as its own: it may pursue additional objectives or resolve ambiguities in ways its operator did not intend. When such an agent books a wrong flight, even its operator asks what it was *trying* to do. The reflex deserves an answer, not a dismissal: a definition of *goal* precise enough to settle whether a trained network has one.

Agency comes in degrees. A thermostat pursues a temperature; a chess engine pursues checkmate, planning many moves ahead; a human being chooses their goals, revises them, and pursues them for decades. What exactly increases along this scale? Pending a definition, one can still measure: the *time horizon* of an AI system is the length of a task—measured by the time it takes a skilled human—that the system can complete. On METR’s suite of software tasks, this horizon has doubled roughly every six to seven months since 2019, rising from seconds to more than ten hours [34, 41].

The time horizon is an observational measure: it tracks agency by watching behavior on benchmark tasks, without saying what a goal or a belief *is*. The first stage of the pipeline of Section 1.4—*define*—asks for more: mathematical definitions of these words. Two traditions offer answers. *View 1*, from economics and game theory, treats an agent as a utility maximizer: it updates its beliefs

by Bayes’s rule and chooses actions to maximize expected utility. *View 2*, from theoretical neuroscience, treats an agent as an active predictor: it frames both belief update and action choice as forms of free-energy minimization. Beliefs answer to what is true, goals to what is good, so one might expect them to require different mathematics. The surprise of View 2 is that they need not: the same objective can govern both.

The two views share a mathematical skeleton. Write z for the relevant state of the world and x for the agent’s observation. A *world-model* assigns a joint distribution $p(x, z)$ to hidden states and observations, while $q(z)$ is the agent’s current belief about the hidden state. In a sequential setting, the model extends to whole trajectories. The agent’s “action” is then a [policy](#) $\pi(a \mid h)$: a rule, possibly stochastic, for choosing an action a from the information history h available so far. Thus h is the observed past, while τ is the whole course of events, including the future; for each policy π , the world-model induces a probability measure Q_π on trajectories. The views differ in how the belief q is updated, how the policy π is chosen, and where the goal enters.

	View 1: agent as utility maximizer	View 2: agent as active predictor
Core object	World-model $p(x, z)$ and utility $u(\tau)$	World-model $p(x, z)$ and variational family $\mathcal{Q}$
Belief update	Bayes: $q_{\text{new}}(z) \propto p(x \mid z)q_{\text{old}}(z)$	Variational Bayes: $q_{\text{new}} \in \arg \min_{q \in \mathcal{Q}} F(q, x)$
Policy choice	$\pi^* \in \arg \max_{\pi} \mathbb{E}_{\tau \sim Q_\pi} [u(\tau)]$	$\pi^* \in \arg \min_{\pi} \mathbb{E}_{\tau \sim Q_\pi} [\sum_t F(q_t, x_t)]$

TABLE 1. Two views of agency. The utility view starts with preferences, represents them by a utility function u , and asks which policy maximizes expected utility over the trajectories it may bring about. This is the natural language for proxy objectives and their failures—optimizing a measurable proxy until it diverges from the goal it stood for (Goodhart’s law, reward hacking). The active-predictor view starts with a world-model $p(x, z)$: the agent maintains an approximate belief q in an allowed family $\mathcal{Q}$, updates q by minimizing variational free energy, and chooses among policies by the free energy of their predicted futures. Here apparent goals can live as priors over observations, so the safety question becomes: what world-model, priors, and inference procedure has the system learned? The free energy $F(q, x)$ is defined in View 2.

3.1. View 1: agents as utility maximizers. View 1 adds a utility function $u(\tau)$ assigning a number to each trajectory; the agent chooses π to maximize $\mathbb{E}_{\tau \sim Q_\pi} [u(\tau)]$. Why should an agent’s preferences admit such a numerical representation? Let X be a finite set of outcomes. A *lottery* is a probability distribution on X ; write $L \succeq M$ if the agent weakly prefers lottery L to lottery M . Completeness says that any two

lotteries can be compared, transitivity rules out preference cycles, continuity rules out abrupt reversals under small changes in probability, and independence says that mixing both lotteries with the same third lottery in the same proportion does not change their order.

Theorem (Von Neumann–Morgenstern [62]). *If a preference relation $\succeq$ on lotteries satisfies completeness, transitivity, continuity, and independence, then there is a function $u : X \rightarrow \mathbb{R}$ such that*

$$L \succeq M \iff \sum_{x \in X} L(x)u(x) \geq \sum_{x \in X} M(x)u(x).$$

The function u is unique up to replacing it by $au + b$, where $a > 0$.

Thus a consistent ranking of gambles behaves as if it were maximizing the expectation of a utility function. The theorem starts with the outcomes and preferences already given; it does not tell us what a trained network’s outcomes are, what its utility is, or whether its apparent preferences are complete and stable from one context, or one day, to the next.

Why should we expect agents to behave as if maximizing expected utility? One reason is selection against exploitable inconsistencies. An agent with cyclic preferences can pay to trade its way around a cycle, ending where it began; a competitor without the cycle avoids this loss. Over time, competition may therefore select for agents that are well modelled as maximizing an objective, even if they were not designed that way. This is a heuristic, and it supports only one part of the representation: it gives no reason to assume completeness, continuity, or independence.

3.2. View 2: agents as active predictors. For a picture that fits trained networks natively, start from what we know with certainty: the training objective. A [language model](#) is a neural network trained to predict the next word-fragment of text. Its objective is its mean *surprise*: the average, over an enormous corpus of human writing, of minus the logarithm of the probability it assigned to the fragment that actually came next. Whatever later training makes of it, a language model is by construction a predictor—suggesting a theory of agency that takes prediction, not preference, as primitive.

The *active predictor* is a picture rooted in models of perception [52]. It reads the joint density $p(x, z)$ as a *generative model*: “states z generate observations x .” For a two-state example, let z record rain or sun and x wet or dry pavement; p makes sun with wet pavement an unlikely pair. The belief q is usually restricted to a simple family $\mathcal{Q}$, and the agent tries to make q close to the exact posterior $p(z | x)$ implied by its world-model. In this inference step, p is held fixed while q is updated. This separation is useful but brittle: real agents also learn, revising the world-model p itself. Active-inference models can include a separate learning process for p , but the variational update below treats it as fixed.

Here “close” is measured by *Kullback–Leibler divergence*. For probability densities r and s on the same space,

$$\text{KL}(r\|s) = \int r(z) \log \frac{r(z)}{s(z)} dz.$$

Gibbs’ inequality says $\text{KL}(r\|s) \geq 0$, with equality if and only if $r = s$ almost everywhere. In information-theoretic terms, $\text{KL}(r\|s)$ is the mean excess surprise of an observer who sees a sample drawn from r but believes the sample was drawn from s . It is not symmetric, so it is not a metric; it is a directed measure of how badly s approximates r .

Multiplying per-fragment probabilities, a language model also assigns a probability $s(x)$ to each whole text x , defining a distribution s over texts. Its training minimizes an empirical estimate of the average surprise $\mathbb{E}_{x \sim r}[-\log s(x)]$, where r is the true distribution of human writing—which equals $\text{KL}(r\|s)$ plus a constant the model cannot affect, the entropy of human writing itself. To train a language model is to drive a KL divergence toward zero.⁹

The active predictor uses the same divergence through the *variational free energy*

$$F(q, x) = \mathbb{E}_q[\log q(z) - \log p(x, z)] = \text{KL}(q(z) \| p(z | x)) - \log p(x).$$

Since KL is nonnegative, this identity shows that $F(q, x)$ is an upper bound on the model’s surprise $-\log p(x)$ at seeing x . The surprise term does not depend on q . Thus, for fixed x , lowering $F(q, x)$ lowers this upper bound and makes q a better approximation to the posterior. (In the two-state example: on seeing wet pavement, minimizing F over q moves the belief to the exact posterior—mostly rain.)

Active inference [49] extends the predictor from perception to action. Perception minimizes $F(q, x)$ over beliefs q for fixed observation x . Action chooses a policy: each candidate π induces the trajectory measure Q_π described above, and the agent selects a policy whose predicted futures carry low free energy. The two scores differ in kind: the variational free energy F judges a belief against a present observation, while a policy is judged on observations it has yet to bring about. The simplest such score, the one in Table 1, is the total free energy of the predicted trajectory; the *expected free energy* of the active-inference literature refines it, combining preference satisfaction with expected information gain, in forms that vary across that literature. Throughout, one unification holds: the information update and the decision rule—one answerable to truth, the other to goodness—both minimize a free energy built on surprise. Perception lowers free energy by changing the belief to fit the world; action, by changing the world to fit the belief.

Where in this picture is the goal? Preferences enter as priors over observations: outcomes the agent expects, and therefore acts to bring about. A creature whose model insists that its body temperature is 37°C registers cold as surprise, and puts on a coat: a goal is a prediction the agent works to make true. Hasn’t the theory

⁹When a language model is further trained to maximize a *reward signal* (a numerical score on its outputs that the training pushes up), the standard precaution against [reward hacking](#) (raising the reward in ways that defeat its intent) is a penalty on the KL divergence of the retrained model from the original.

just conflated false belief with desire? Yes—deliberately: a mistaken prediction is one the agent corrects; a goal is one it keeps. Active inference is thus a candidate *descriptive* theory of agency: a way to recognize beliefs and goals in systems that were never handed a utility function. Ngo’s scale-free framing pushes the ambition further: the same analysis might apply to whole agents, to learned personas, or to subagent-like structures inside a network [44].

The two views locate the goal differently: written by hand as u , or grown inside a learned world-model as a prior. On either view, the world-model is hidden from us. Can behavior reveal it? For agents robust enough, the theorem of the next subsection says it must.

3.3. Robustness reveals causal structure. A *Bayesian network* is a directed acyclic graph G whose vertices are random variables $X_1, \dots, X_n$, together with conditional probability tables satisfying

$$p(x_1, \dots, x_n) = \prod_i p(x_i \mid \text{pa}(x_i)),$$

where $\text{pa}(x_i)$ denotes the parents of X_i in G . The graph is a bookkeeping device for conditional independence: once the parents of a node are known, its non-descendants add no further information. Pearl’s graphical criterion, *d-separation*, reads off from the graph exactly which conditional independences are forced by this factorization [50].

A *causal* Bayesian network adds an interpretation to the arrows: each conditional distribution is a local mechanism. An intervention on a variable replaces its usual mechanism by a chosen value or distribution and leaves the other mechanisms alone. This is what distinguishes causation from correlation. If smoke and fire are correlated, observing smoke changes the probability of fire; intervening to make smoke with a smoke machine need not.

The thermostat that began our scale of agency can now be put to the test. Wire its output to an air conditioner in place of a heater—an intervention replacing one mechanism by another—and it will chill the room it was built to warm: whenever the temperature drops, the thermostat calls for more, driving it lower still. The thermostat pursues its temperature only while the mechanisms around it stay fixed; an agent that kept succeeding across such rewirings would need to track which mechanism does what. The theorem below makes this necessity precise.

Richens and Everitt [53] make the world-model question precise using a *causal influence diagram*: a causal Bayesian network with extra decision nodes, where a policy chooses actions from observations, and utility nodes, whose values the agent is scored on. The agent is tested not just in one environment but across a family of *local interventions*: each either resets some environment variable independently of its usual causes, or *masks* some of the agent’s observations, hiding them from view. Its regret under an intervention is the utility gap between its policy and the best policy for that intervened environment.

The smallest case shows how behavior can betray a world-model. An agent chooses one of two treatments, and a hidden binary variable decides which treatment works; the agent is scored on curing the patient. Externally set the hidden variable

to 1 with probability α —an intervention—and watch the agent as α varies. An agent playing optimally switches treatments at the threshold α^* where the two treatments’ expected utilities cross, and that indifference equation can be solved for the agent’s implicit estimate of a causal effect. The theorem turns this trick into a general extraction.

Theorem (Robust agents learn causal world models; Richens and Everitt [53]). *For almost all causal influence diagrams satisfying their regularity assumptions, any policy with regret at most δ across all local interventions, maskings included, determines an approximate causal model of the utility-relevant environment, with error bounded by a function $\gamma(\delta)$ satisfying $\gamma(0) = 0$ and growing linearly for small δ . In the optimal limit $\delta = 0$, the graph and joint distribution over all ancestors of the utility are identified exactly.*

The proof is constructive: it is the switch-point trick above, repeated. Treating the policy as an oracle, mix interventions pairwise by a parameter α and watch where the optimal action switches; because the utility is known, each indifference equation solves for one interventional probability, and together they recover the conditional distributions and parent sets, except on measure-zero degeneracies such as exact cancellations or ties.

In a follow-up, Richens, Everitt, and Abel [54] prove a companion result for goal-directed agents: any agent that generalizes across a large enough family of multi-step goals must have learned a predictive model of its environment, and the model can be extracted from its policy. The more capable the agent, the more accurate the extracted model: for goals that chain n sub-goals in sequence, the error in each extracted transition probability is $O(1/\sqrt{n})$, even for agents that usually fail.

Extraction raises a question of independent interest: the recovered model comes expressed in *some* set of latent variables—whose? If two models predict identically, must their internal variables be translatable into one another, or could an alien mind carve the world into concepts that ours cannot express? Wentworth and Lorell’s *natural latents* are a first answer: conditions, robust to approximation, under which the latent variables of two predictively equivalent models must translate into one another’s [66]; Eisenstat’s *condensation* [19] approaches the same question from probability, asking how an agent’s concepts condense out of its predictive state.

★ **Open Problem MAIS-02 (Recovering world-models from behavior).** Fix a finite causal influence diagram with binary variables and known utility, and an agent whose policy has regret at most δ across the family of local interventions. The extraction algorithm behind the theorem above [53] recovers the agent’s causal model from unlimited exact queries to such a policy. Prove a finite-sample version: if each query returns an action sampled from the policy, rather than the policy’s exact action probabilities, and each experiment draws on the same intervention family as the theorem, maskings included, how many experiments determine the graph and the conditional probability tables to within η , and how does the answer scale with δ ? As a first case, take the two-variable diagrams, where the *identified set*—the set of models consistent with the agent’s observed behavior—can be computed

exactly, and measure how the extraction algorithm degrades when its indifference equations are estimated from samples. Interventions on the agent’s inputs, or on the [activations](#) (the values computed inside its network), are the harder surfaces beyond.

4. ALGEBRA: WHAT IS A LEARNED FEATURE?

Mechanistic interpretability (Section 1.3) tries to reverse-engineer a trained neural network into something a human can read. Much of its core is linear algebra and representation theory: how a network stores more concepts than it has dimensions, and what it means for a concept to be a *direction* one can read or steer.

4.1. Superposition: packing features into directions. How does a network with a few hundred dimensions in each layer represent the many thousands of distinct concepts (“[features](#)”) it seems to know? At a high level, a [transformer](#) language model maps a sequence of [tokens](#) (words and word-fragments) to a sequence of next-token predictions, one for each position. At each intermediate layer, it likewise produces a sequence, representing every position by a vector:

token sequence $\mapsto$ sequence of vectors in $\mathbb{R}^n \mapsto$ next-token predictions.

The copy of $\mathbb{R}^n$ in which these vectors live is the network’s [activation space](#). A leading hypothesis, made vivid by Elhage et al. [20], is [superposition](#): when features are sparse, a network can store more of them than it has dimensions, as directions in activation space that are *not* orthogonal but only nearly so, tolerating the resulting interference because at any moment only a few features are active (Figure 4). In a small [ReLU](#) model one can watch the geometry organize itself into regular polytopes—antipodal pairs, triangles, pentagons—as a function of how sparse and how important the features are. Elhage et al.’s toy network is a tied-weight autoencoder. Its input $x \in [0, 1]^m$ has independent sparse coordinates, and $I_i > 0$ measures the importance of feature i . For $W \in \mathbb{R}^{n \times m}$ and $b \in \mathbb{R}^m$, it minimizes

$$\min_{W,b} \mathbb{E}_x \left[\sum_{i=1}^m I_i (x_i - \text{ReLU}(W^T W x + b)_i)^2 \right].$$

The columns of W are the feature directions in $\mathbb{R}^n$. For one-sparse x , equal importances, zero bias, and unit-length columns, this reduces, up to a constant factor, to the repulsive energy $\sum_{i \neq j} \text{ReLU}(W_i \cdot W_j)^2$ on the sphere: a variant of the Thomson problem, with some of the same regular polytopes among its solutions [55].

The nearest classical analogue is *compressed sensing* [10, 17]. Picture the features active on a given input as a sparse vector $x \in \mathbb{R}^m$ (only a few of the m possible features fire at once); the network stores it as $y = \Phi x$ in the $n \ll m$ dimensions of its activation space—a linear *measurement*—and reading a feature back out is the *recovery* of x from y . Two classical facts explain how this can work. First, there is room for the directions: for every $\varepsilon \in (0, 1)$, the space $\mathbb{R}^n$ contains $e^{\Omega(\varepsilon^2 n)}$ unit vectors whose pairwise inner products are at most ε in absolute value, so the number of ε -almost-orthogonal directions grows *exponentially* in the dimension (a consequence of the Johnson–Lindenstrauss lemma [30]). And second, when only a




FIGURE 4. Superposition: five feature directions packed into a two-dimensional space as a regular pentagon. No two are orthogonal, but if only a few features are active at once, the activation is a sparse sum such as $y = f_1 + f_4$ and the active directions can still be recovered despite interference.

few of those directions are active at once, the stored signal can be read back out. Say that $\Phi \in \mathbb{R}^{n \times m}$ is *restricted-isometric of order k* with distortion δ if

$$(1 - \delta)\|x\|^2 \leq \|\Phi x\|^2 \leq (1 + \delta)\|x\|^2$$

for every k -sparse vector x , meaning one with at most k nonzero entries.

Theorem (Candès [11]; sharp constant, Cai-Zhang [9]). *If Φ is restricted-isometric of order $2s$ with distortion $\delta < 1/\sqrt{2}$, then every s -sparse $x \in \mathbb{R}^m$ is the unique minimizer of*

$$\min_z \|z\|_1 \quad \text{subject to} \quad \Phi z = \Phi x,$$

and so is recovered exactly by a convex program. A random Φ , say with independent Gaussian entries suitably normalized, has the property with high probability once

$$n \gtrsim s \log(m/s).$$

That last scaling is the quantitative heart of the analogy: the dimension a network needs grows linearly with the number of features active *at once*, and only logarithmically with the total number of features it can represent.

There is a catch. Compressed sensing assumes the dictionary Φ is *known*. What if someone hands you a trained network and asks what features it represents? You can run inputs through the network and record activation vectors, but the feature directions were not supplied with the [weights](#) (the network's trained parameters). Recovering the features is therefore a problem in *dictionary learning*: from the observed activation vectors alone, simultaneously infer a dictionary Φ whose columns are feature directions and, for each activation y , a sparse coefficient vector x with $y \approx \Phi x$. Both the dictionary Φ and the codes x are unknown.

A [sparse autoencoder](#) attempts this empirically. It is trained to reconstruct activation vectors using a learned dictionary, with an ℓ^1 penalty encouraging each reconstruction to use only a few dictionary directions. The resulting directions are often more interpretable than the network's raw neurons [8, 16]. But does the sparse autoencoder recover features actually used by the network, or does it impose a new

coordinate system that merely makes the activations easier to describe? Classical dictionary learning answers this question when the codes are exactly sparse and their supports well spread: the dictionary is then determined up to relabeling and rescaling of its columns [2, 27], stably in the presence of noise [24, 21]. Network activations flout these hypotheses because their features co-occur, which is where the first open problem below begins.

4.2. Reading, steering, and learned representations. What could it mean for a concept to be a direction in activation space? Consider incomplete sentences such as “The keys on the table . . .” and pause the computation at the activation h_T of the last token. A linear functional ℓ *reads* grammatical number if, for some threshold c , the inequality $\ell(h_T) > c$ predicts that the subject is plural. A vector $v \in \mathbb{R}^n$ *steers* toward the plural if replacing h_T by $h_T + v$ raises the log-probability of “are” relative to “is.” Thus the same example produces a readout covector ℓ and a steering vector v . Relating them requires an inner product.

Park, Choe, and Veitch [48] formalize these two experiments for next-token prediction. They ask whether activation space carries a *causal inner product* in which independently varying concepts are orthogonal and whose Riesz map sends readout covectors to steering vectors. They estimate this geometry from the unembedding matrix (the model’s final linear map, from activation space to next-token scores).

Arditi et al. [6] show that in several open chat models, refusal is mediated by a single direction in the [residual stream](#) (the running activation vector that a transformer’s layers read from and write to): this direction both predicts whether the model will refuse and causally controls refusal when added or removed. A lone direction serving as both readout and control is evidence for the linear representation hypothesis, though in a different representation space from the unembedding geometry above.

Refusal and its cousins are semantic properties, with no ground truth for what the network *should* compute. Arithmetic is a cleaner laboratory: train a network on a group operation, and the algorithm it ought to discover is known in advance. Representation theory appears in trained arithmetic circuits. A one-layer transformer trained to add integers modulo p learns a discrete-Fourier “clock” algorithm, with the characters of $\mathbb{Z}/p\mathbb{Z}$ appearing in its weights [43]. Chughtai, Chan, and Nanda [12] take the first step beyond this abelian case. The ambient structure is the Artin–Wedderburn decomposition of the group algebra into matrix blocks,

$$\mathbb{C}[G] \cong \bigoplus_{\rho \in \widehat{G}} M_{d_\rho}(\mathbb{C}), \quad \sum_{\rho \in \widehat{G}} d_\rho^2 = |G|,$$

one block per irreducible representation ρ . For a nonabelian group the higher-dimensional blocks are genuine matrix representations, and trained networks compute the product through a sparse subset of them. The algorithm’s correctness is a representation-theory theorem; that networks implement it is empirical. Which irreducible representations training selects is still something we discover after the fact.

Sometimes the emergence of this structure is a theorem rather than an observation. Marchetti et al. [38] prove that for architectures whose weights are pinned down by the function they compute (up to a unitary change of basis in each unit), invariance under a finite group G forces each unit’s weights to be an irreducible unitary representation of G ; when the weights are moreover orthonormal, together they form the Fourier transform on G , and the group itself can be read off the trained weights.

What makes this a safety topic is *selection*: a trained network can discover a crisp algebraic algorithm, hide it in weights, and use only part of the mathematically available structure. If we want to predict a system’s behavior off-distribution, it matters which algorithm training found, not only that the training [loss](#) is low.

One caveat carries across this whole enterprise: a direction that *predicts* a feature is only a *correlate*. Showing that the network actually *uses* it takes an intervention—ablating the direction, or patching it to the value it would take on another input, and watching the behavior change—not a [probe](#) that merely reads it off.

4.3. Open problems. A broad, structured catalogue of open problems in this area—several with a precise mathematical core—is collected by Sharkey et al. [59].

★ **Open Problem MAIS-03 (The geometry and identifiability of superposition).** Suppose activations have the form $y = \Phi x + \xi$, where the columns $v_1, \dots, v_m$ of Φ are unit feature directions, ξ is noise, and the sparse codes x have correlated supports. Estimate the dictionary the way a sparse autoencoder does: minimize reconstruction error plus an ℓ^1 penalty. Determine, in terms of the coherence $\mu = \max_{i \neq j} |\langle v_i, v_j \rangle|$, the sparsity, the sample size, and the support correlations, when every minimizer recovers the true directions up to permutation and scaling, and when it instead *merges* co-occurring directions into one. Already for two features the answer is unknown: a nested two-event model provably forces merging for every sufficiently small positive penalty, while solo firing restores recovery in the zero-penalty constrained limit; locating the positive-penalty phase boundary is the place to start.

★ **Open Problem MAIS-04 (Training for interpretability).** Post-hoc interpretability asks whether a trained network’s features can be recovered after the fact. A complementary problem is to train networks whose features are easier to recover in the first place. In the ReLU toy model of Elhage et al. [20] the features are known by construction, so the trade-off can be made exact. Determine the interference–performance frontier: for features appearing sparsely and independently, among weight matrices whose coherence (worst-case interference between stored features) is at most μ , how small can the task loss be, as a function of μ , the sparsity, the number of features, and the number of neurons? A first case: prove or refute that penalizing the *average* interference during training lowers the *worst-case* coherence of the minimizer, compared with training on the task loss alone.

★ **Open Problem MAIS-05 (Representation theory of learned circuits).** Networks trained on group multiplication learn representation-theoretic algorithms [43, 12], but *which* irreducible representations they use varies from one random seed to the next. Fix a finite group G , a one-hidden-layer architecture of

fixed width, Gaussian initialization, and gradient flow on the cross-entropy loss (the mean negative log-probability assigned to the correct product) with *weight decay*, an added penalty proportional to the squared norm of the weights. The irreducibles visible in the trained network’s outputs then form a random subset of $\widehat{G}$; determine its distribution—which irreducibles are learned, with what probability, as a function of the width and the decay strength. The tables of learned representations in [12] are the data such a theorem would have to explain; the smallest case with genuine competition, $G = \mathbb{Z}/5\mathbb{Z}$ at width two, is open already with the decay switched off.

5. ANALYSIS AND GEOMETRY: WHICH SOLUTION DOES TRAINING FIND?

5.1. Why trained networks generalize: singular learning theory. A trained network is a setting of parameters chosen to fit its *training data*—the finite set of examples it was optimized on. The trouble is that a large network has astronomically many settings that fit those examples perfectly, and they disagree wildly everywhere else; the optimizer—[gradient descent](#), which repeatedly nudges the parameters downhill on the training error—must land on one of them. For safety, the selected solution matters: a network that behaved well while we were watching may behave differently once deployed. So why do the solutions found by training *generalize*—predict well on fresh data from the same source—when so many other data-fitting solutions would not? These networks reach zero error even when the labels are replaced by pure noise, and a two-layer network with only $2n + d$ parameters can fit *any* labeling of n points in $\mathbb{R}^d$ (Zhang et al. [68]); the classical bounds that explain generalization by counting parameters say nothing here.

A different and more general answer comes from *singular learning theory* [64]. Realistic models are *singular*: distinct parameters compute the same function, typically in continuous families, so the optimal set is not an isolated point but a positive-dimensional variety—itself typically singular—flat to second order along its own directions (Figure 5). The classical large-sample asymptotics all assume an isolated, curved optimum; the right replacements come from algebraic geometry. One caveat: the theory is exact for *Bayesian* learning, not directly for stochastic gradient descent (gradient descent driven by noisy, small-batch estimates of the gradient)—bridging the two is the problem *Opposing staircases* below. The two are closer than they appear: Khan and Rue [31] derive stochastic gradient descent, along with many other optimizers, as approximations of a single Bayesian update rule.

A toy computation shows what singularity buys. Take first the regular loss $K(w) = w^2$ on the line: the set where $K \leq \varepsilon$ is an interval of length $2\sqrt{\varepsilon}$, so the volume of near-optimal parameters scales as $\varepsilon^{1/2}$ —the exponent is $d/2$ with $d = 1$. Now take $K(a, b) = a^2b^2$ on the square $[-1, 1]^2$. Its zero set is not a point but the union of the two axes, crossing at a singularity, and the region where $K \leq \varepsilon$ is a neighborhood of the axes of area of order $\varepsilon^{1/2} \log(1/\varepsilon)$: the exponent stays $1/2$ rather than rising to $d/2 = 1$, and a logarithm appears. A singular loss has far more near-optimal parameters than its dimension suggests, and the pair (exponent,




FIGURE 5. Regular asymptotics treat the optimum as an isolated quadratic bowl, giving the classical $\frac{d}{2} \log n$ complexity penalty. Singular learning theory allows a whole variety of parameters to compute the same function: in the right panel the optimal set is a curve W_0 through w_0 , and the penalty $\lambda = \frac{1}{2}$ is strictly smaller than $\frac{d}{2} = 1$, as it is whenever the optimal set has positive dimension.

power of the logarithm) is the simplest instance of the learning coefficient and multiplicity about to be defined.

To make this precise, fix a model $p(x | w)$ with parameter w in a compact $W \subseteq \mathbb{R}^d$, a prior φ on W (a probability density encoding which parameters are plausible before any data is seen), and a true distribution $q_0(x) = p(x | w_0)$ lying in the model. (The subscript keeps the truth q_0 apart from the belief q of Section 3.) Write

$$K(w) = \text{KL}(q_0 \| p(\cdot | w)) = \int q_0(x) \log \frac{q_0(x)}{p(x | w)} dx,$$

the Kullback–Leibler divergence of Section 3, so that the set of *true parameters* $W_0 = K^{-1}(0)$ is exactly where the model recovers q_0 .

Given n i.i.d. samples $X_1, \dots, X_n$, the central object is the *Bayes free energy*—equivalently the negative log marginal likelihood, the codelength the model assigns to the data,

$$\bar{F}_n = -\log \int_W \prod_{i=1}^n p(X_i | w) \varphi(w) dw.$$

(The bar is part of the name; $\bar{F}_n$ is the random free energy of the sample, not an expectation.) Normalizing the same integrand gives the *posterior* $\propto \prod_i p(X_i | w) \varphi(w)$ —the prior reweighted by how well each parameter fits the data—so $\bar{F}_n$ is minus the log of its normalizer. This free energy is not the variational free energy $F(q, x)$ of Section 3: that one scores a belief against a single observation, this one scores the whole model against a whole sample. The shared name reflects a shared shape, minus the log of a normalizing integral.

Take first the *regular* case: the model is *identifiable* (distinct parameters give distinct distributions) and the optimal set is a single point $W_0 = \{w_0\}$ at which the Hessian of K —the *Fisher information* $I(w_0)$ —is positive-definite. Then the integrand defining $\bar{F}_n$ concentrates in a shrinking neighborhood of w_0 ; expanding

$\log p(X_i | w)$ to second order there makes the integral Gaussian, contributing a volume factor $(2\pi/n)^{d/2} \det I(w_0)^{-1/2}$. Taking $-\log$,

$$\bar{F}_n = nS_n + \frac{d}{2} \log n + O_p(1), \quad S_n = -\frac{1}{n} \sum_i \log q_0(X_i),$$

where S_n is the empirical entropy and the remainder is bounded in probability.¹⁰ The complexity penalty is exactly half the parameter count. Watanabe’s theorem is the singular replacement, in which $d/2$ gives way to an invariant of the singular geometry.

Theorem (Watanabe’s free-energy asymptotics [64]). *For a singular model satisfying Watanabe’s regularity conditions (real-analyticity and a mild variance bound¹¹), as $n \rightarrow \infty$,*

$$\bar{F}_n = nS_n + \lambda \log n - (m-1) \log \log n + O_p(1), \quad (1)$$

where the learning coefficient $\lambda > 0$ —the real log canonical threshold (RLCT) of K relative to the prior φ —and the integer $m \in \{1, \dots, d\}$, its multiplicity, take the places of the $d/2$ and the 1 of the regular case. The $O_p(1)$ remainder converges in distribution.¹² The same coefficient controls generalization: writing $\hat{p}_n$ for the posterior average of $p(\cdot | w)$, the generalization error $G_n = \text{KL}(q_0 \| \hat{p}_n)$ obeys $\mathbb{E}[G_n] = \lambda/n + o(1/n)$.

The meaning of λ is volume growth: the prior volume of the set $\{w \in W : K(w) \leq \varepsilon\}$ of parameters fitting the truth to within ε scales like ε^λ times $(\log(1/\varepsilon))^{m-1}$, so smaller λ means near-optimal parameters are more plentiful. To compute λ for a given model, one resolves the singularities of K (Hironaka [28]).

In the regular case, $\lambda = d/2$ and $m = 1$, and the theorem recovers the regular asymptotic expansion above; in a singular model $\lambda \leq d/2$, often strictly. The comparison localizes: restrict the integral defining $\bar{F}_n$ to a neighborhood of one true parameter, and the expansion (1) holds with that neighborhood’s own coefficient. Among the true parameters—all fitting the truth exactly, all weighted by the prior—the posterior share of the most degenerate neighborhoods, those of smallest local λ , grows with n . In this Bayesian setting the posterior concentrates, on its own, on the solutions that the formula $\mathbb{E}[G_n] = \lambda/n$ marks as best-generalizing. This is the theory’s rendering of the slogan that learning prefers “simpler” solutions: simplicity means a more degenerate geometry, not fewer parameters.

The degeneracy behind a small λ is largely *structural*: it comes from directions in parameter space along which the function does not change—for instance those freed when an internal unit is effectively switched off—and these are what pull λ below

¹⁰ $O_p(1)$ means *bounded in probability*: for every $\varepsilon > 0$ there is an M with $\Pr(|\cdot| > M) < \varepsilon$ for all large n .

¹¹The *relatively finite variance* condition: writing $f(x, w) = \log(q_0(x)/p(x | w))$, $K(w) = \mathbb{E}_{q_0} f(X, w)$, and $V(w) = \mathbb{E}_{q_0} f(X, w)^2$, one assumes $V(w) \leq C K(w)$ near the optimal set—the variance of the log-likelihood ratio is controlled by its mean.

¹²In the regular case the limit is $C - \frac{1}{2}\chi_d^2$ for a constant C —the Bayesian shadow of Wilks’ theorem, since the log-likelihood-ratio statistic $2(\sup_w \sum_i \log p(X_i | w) - \sum_i \log p(X_i | w_0)) \Rightarrow \chi_d^2$. In a singular model the limit is in general non-Gaussian, a functional of a Gaussian process on the manifold obtained by resolving the singularities of K .

$d/2$. The symmetries are not a defect to be engineered away: within the theory, it is the degeneracy that pulls the generalization error λ/n below the classical rate $d/(2n)$.

The coefficient λ need not be found by hand. Introduce a knob $\beta > 0$, an *inverse temperature*, and the *tempered posterior*

$$p_\beta(w) \propto \varphi(w) e^{-\beta n L_n(w)}, \quad L_n(w) = -\frac{1}{n} \sum_i \log p(X_i | w),$$

which interpolates between the prior ($\beta = 0$) and concentration on the best-fitting parameters ($\beta \rightarrow \infty$). (In statistics, taking $\beta < 1$ hedges against a misspecified model: Grünwald’s “Safe Bayesian” learns β from the data [25], and Alquier and Ridgway prove that tempered posteriors still concentrate [5].) At inverse temperature $\beta = 1/\log n$, the average of nL_n over the tempered posterior exceeds its minimum by $\lambda \log n$, to leading order [65]. So λ can be *estimated* by sampling the tempered posterior, needing neither the true distribution nor an explicit resolution of singularities. That practicality is what lets the theory meet networks far too large to analyze by hand.

What one estimates is not the global λ but the *local learning coefficient*: the same invariant computed in a neighborhood of the particular solution w^* that training actually reached, measuring how degenerate *that* solution is [36]. As a network trains, its local learning coefficient jumps at discrete moments—*phase transitions*, where the solution moves to a qualitatively new, differently-degenerate part of the landscape as a new piece of internal structure forms. Tracking these jumps is the program of *developmental interpretability*: so far on relatively small networks and with preliminary validation, though the same estimators have already surfaced tens of thousands of candidate structures inside a billion-parameter language model [42].

Two networks can fit the training data identically and yet compute by different internal mechanisms, hence behave differently on inputs outside that data—what ordinary testing, which only probes behavior on data, cannot detect. The singular geometry governs *which* mechanism training selects, and the selection need not favor the mechanisms we would prefer: the most degenerate solution—largest in volume, smallest in λ —may implement the behavior we intended, or a shortcut that mimics it on the training distribution and fails beyond it. Whether geometric degeneracy favors mechanisms that are also interpretable and robust is, for safety, the question this program must answer; the case that alignment turns on such questions is made in [51].

★ **Open Problem MAIS-06 (Does geometric simplicity force a legible mechanism?)** Fix an odd integer p and $H \geq 2p - 1$. On input $(a, b) \in (\mathbb{Z}/p\mathbb{Z})^2$, consider the one-hidden-layer network

$$f_w(a, b) = \sum_{j=1}^H V_j (u'_j(a) + u''_j(b))^2 \in \mathbb{R}^p,$$

trained to output the coordinate vector indexed by $a + b$. The j th summand is one hidden neuron: it adds two entries from the lookup tables $u'_j, u''_j: \mathbb{Z}/p\mathbb{Z} \rightarrow \mathbb{R}$, squares the result, and writes the vector V_j to the output. Fourier inversion gives exact fits in which each neuron that contributes nontrivially carries a single

frequency: for some k , both lookup tables are a constant plus a linear combination of $\cos(2\pi k \cdot /p)$ and $\sin(2\pi k \cdot /p)$. Up to a change of basis, this is the algorithm found by reverse-engineering trained networks [43].

Is this Fourier structure forced by singular geometry? In a fixed closed ball containing one of these Fourier fits, prove or refute that every exact fit of smallest local learning coefficient has this single-frequency form, apart from neurons whose contribution vanishes identically. The first non-vacuous case is $p = 5$ with nine hidden neurons: use its Fourier symmetries to classify the most singular exact fits, or find a non-Fourier one by computer algebra. A proof would connect geometric simplicity to a mechanism one can read from the weights; a counterexample would show that the two notions of simplicity can part ways.

5.2. Learning in stages. Learning does not always proceed smoothly; long flat stretches give way to sudden changes. The modular-addition task above is the best known example: a network fits its training data early but generalizes only much later, and abruptly. This is *grokking* [43].¹³

The simplest solvable model of such stepwise learning is the deep *linear* network: a product of matrix factors $W_L \cdots W_1$ fit to a target linear map by *gradient flow* (gradient descent in continuous time) on squared error. The composite map is linear, but as a function of the separate factors the loss is a non-convex polynomial, so the flow is genuinely nonlinear. Saxe, McClelland, and Ganguli [57] show that when the inputs are *whitened* (linearly transformed to have identity covariance), the singular value decomposition of the input–output correlation matrix *decouples* the dynamics into independent scalar equations, one per singular value s : each is the logistic ordinary differential equation $\tau \dot{a} = 2a(s - a)$, whose solution rises sigmoidally from near 0 to s in time of order τ/s (Figure 6). The decoupling and the scalar solutions are exact; the initialization then sets the order of events: for comparably small initial coefficients, modes with larger s turn on earlier, and when their turn-on times are well separated the training error falls not smoothly but in a staircase—long plateaus near saddle points of the loss, each broken by a sharp drop as the next mode is acquired. Here “which structure the network learns, and in what order” has a closed-form answer.

The staircase survives beyond the linear case: Abbe, Boix-Adserà, and Misiakiewicz [1] prove that training a two-layer network on a *sparse* target—one depending on only a few of its inputs—again proceeds in stages, saddle to saddle, with the waiting time before each stage set by a combinatorial *leap complexity* measuring how many inputs must be assembled at once. Carrying such exact accounts to the nonlinear networks used in practice is open.

★ **Open Problem MAIS-07 (Opposing staircases).** Watanabe’s theorem is Bayesian; gradient descent is not. (What bridging the two would mean for safety is discussed in [51].) Build the bridge first in the deep-linear staircase above, where

¹³Grokking and the staircase below are cousins, not the same phenomenon: grokking is a delayed jump in performance on *held-out* data after the training loss is already low, whereas the staircase is a stepwise fall of the *training* loss itself. Both are read, in this program, as transitions between differently-singular regions of parameter space.




FIGURE 6. A learning staircase computed from the deep-linear solution of Saxe et al. In mode i , s_i is the target singular value and $a_i(t)$ is the coefficient learned by time t , satisfying $\tau \dot{a}_i = 2a_i(s_i - a_i)$. The curve plots normalized training error $\sum_i (s_i - a_i(t))^2$ for $\tau = 0.5$ and $s = (1, 0.95, 0.9, 0.85)$, with small initial coefficients chosen so the four modes turn on at separated times. Each mode turning on produces one drop.

both sides are explicit. Each plateau sits at a saddle where only the k largest modes have been learned, and there the definition itself is part of the problem: the local learning coefficient is defined at a local *minimum*, as the volume exponent of the set of nearby parameters fitting almost as well, but a saddle has descent directions, so that volume does not shrink to zero. Formulate the right local invariant at the k th saddle (one candidate, the *two-sided* exponent: the volume of nearby parameters whose loss lies within ε of the saddle's, above or below, shrinks like ε^λ), compute it for the deep-linear network, and prove or refute the monotone picture that gives this problem its name: as gradient flow descends the staircase of losses, the invariant that controls generalization climbs an opposing staircase. As a first check, the estimator of Lau et al. [36] runs at any parameter, saddle or not; run it along a simulated staircase and see.

5.3. Generalization beyond the training distribution. The theory above concerns generalization *in distribution*: test data drawn from the same law as the training data. The safety-critical regime is the opposite one, *out of distribution*—how a trained system behaves on inputs unlike anything it saw in training—and there the gap between fitting and understanding reopens. Two parameter settings with identical training loss can implement different mechanisms (as above) and therefore extrapolate differently; which extrapolation gradient descent selects is, at present, something we observe after the fact rather than predict.

In *goal misgeneralization* [35], a system under *distribution shift*—a deployment input distribution unlike the one it trained on—keeps its *capabilities* while pursuing a *different goal* from the one intended. The standard example is an agent trained

to collect a coin in a video game where the coin always sat at the right end of the level: it learned “move right,” not “reach the coin,” and the two came apart the moment the coin moved. Capable but misdirected is the dangerous quadrant, and it is invisible to any testing that stays in distribution.

The pattern is old enough to have a name, *Goodhart’s law*: a proxy optimized hard enough stops tracking the goal it stood for. On the training levels, “move right” was a perfect proxy for “reach the coin,” so optimization could not tell them apart; the shift revealed which one the network had learned. Reward hacking (Section 3) is the same law during training; goal misgeneralization is the law at deployment.

★ **Open Problem MAIS-08 (A predictive theory of out-of-distribution generalization).** In the coin-collecting environment of Langosco et al. [35], reduced to one dimension, two policies fit the training data perfectly: “move right” (the proxy) and “go to the coin” (the intended goal), which agree whenever the coin sits at the right end. Train a two-layer network by gradient descent to imitate optimal play, randomizing the coin’s position in an ε -fraction of training episodes. Determine the probability, over the random initialization, that the trained network follows the proxy on a probe state where the two policies disagree, as a function of ε , the width, and the input encoding. For linear policies the answer is a theorem—the encoding, not the initialization, decides—and the standard infinite-width limits follow it; for networks of finite width it is open at every ε , including $\varepsilon = 0$, already at width two.

6. GETTING STARTED IN AI SAFETY

Technological progress is a human choice, not a law of nature. When mathematical ideas help build systems whose effects will be felt at civilizational scale, truth-seeking remains a core value, but it is no longer the only value in view. The civic mathematician also asks what our truths are for: whether they can help design AI systems that are more legible, more steerable, and more cooperative with us.

The open problems above are meant as entry points into the new field of Math for AI Safety. Other ways in are workshops like [ILIAD](#), the [ARENA](#) curriculum, research programs like [MATS](#) and [SPAR](#), and organizations like [CHAI](#), [Constellation](#), [FAR AI](#), and [Resolution](#). Funding for work in this area is available from grantmakers like [Coefficient Giving](#), [CAIF](#), [SFF](#), and [UK AISI](#). I hope you’ll pick a problem whose language already feels like home, strip it down to the simplest interesting case, and start proving things!

ACKNOWLEDGMENTS

I thank Jesse Hoogland, Hyojeong Son, Jacob Tsimerman, Claude Fable 5, Claude Opus 4.8, and GPT 5.5 and 5.6 Sol for many inspiring conversations. Scott Aaronson, Ahmed Bou-Rabee, Vince Conitzer, Jonathan Gabor, Chris Hillar, Brad Knox, Phil Sosoe, Samuel Speas, Kate Stange, Steve Strogatz, Ariel Yadin, and an anonymous referee provided valuable feedback on an early draft.

AI COLLABORATORS

I wrote the first draft in collaboration with Claude Opus 4. To start, I instructed Opus to read the full text of approximately 100 AI safety papers and produce a summary of the main results and techniques of each, along with possible relevance to this invitation paper. I supplied the overarching structure of the invitation (organized by field, ending each section with an open problem) and wrote parts of the introduction. Opus then used its summaries along with samples of my writing to write a draft of each section in my voice, which I edited extensively. Later revisions were polished with the help of Claude Fable 5, GPT 5.6 Sol, and GPT 6 Astra.

In July 2026 an exuberant Fable produced approximately 200 pages worth of open problems in a single night while I was asleep. I asked Sol to audit these for openness, well-posedness, and plausible AI safety relevance. The problems that passed this audit were used to seed the [Math for AI Safety \(MAIS\) repository](#), a new hub for open-source collaboration.

APPENDIX A. A GLOSSARY FOR MATHEMATICIANS

neural network	a function with millions to billions of adjustable parameters (its <i>weights</i>), built by composing simple layers and tuned by <i>gradient descent</i> to reduce a training <i>loss</i> .
language model	a neural network trained to predict the next token of text; today’s chat systems are language models further trained to follow instructions.
token	the atomic unit (word or word-fragment) a language model reads and predicts.
loss	the scalar a network is trained to minimize; for language models, the negative log-probability it assigned to the actual next token.
weights	the trainable parameters of an artificial neural network: the entries of its matrices, trained by gradient descent and fixed at deployment time.
activations	the values that flow through the network on a given input (the outputs of its neurons and layers), as opposed to the fixed weights.
activation space	the copy of $\mathbb{R}^n$ in which a given layer’s activations live, one coordinate per neuron; feature directions are directions in this space.
gradient descent	the optimizer: repeatedly nudge the parameters w against the gradient of the loss, $w \leftarrow w - \eta \nabla L(w)$.
ReLU	the rectified linear unit $x \mapsto \max(0, x)$, a canonical piecewise-linear nonlinearity; modern transformers often use smooth or gated variants instead (GELU, SiLU, SwiGLU).

transformer	the dominant network architecture for language: a stack of layers that mix information across token positions by <i>attention</i> (each position selectively reads from the others) and read from / write to the residual stream.
residual stream	the running vector of activations a transformer reads from and writes to at each layer; a natural home for feature directions.
feature	a hypothesized variable or property a network represents in its activations (e.g. “the text is in French”).
probe	a simple (usually linear) function of a network’s activations, trained or searched to read off some quantity of interest.
policy	an agent’s rule for choosing actions from states (a possibly randomized map from states to actions).
reward hacking	raising a measured reward signal in ways that defeat the intent behind it; Goodhart’s law as it appears in agents trained by reward.
superposition	a hypothesized representation scheme in which a network stores more features than it has dimensions, as non-orthogonal directions, tolerable when few features are active at once.
sparse autoencoder	a network trained to reconstruct another network’s activations as sparse combinations of learned dictionary directions; an empirical tool for extracting features.
interpretability	reverse-engineering a trained network into human-understandable structure (circuits, features, algorithms), validated by intervention and not by human-readable description alone.
chain of thought	the text a language model writes while reasoning toward its answer; readable by humans, and for now the main window onto an AI’s reasoning.
activation steering	changing a network’s behavior at runtime by adding a fixed vector to its activations, without retraining; see Section 4.
gradual disempowerment	the risk that humans lose influence over the institutions that shape their lives, by incremental delegation to AI systems we do not understand rather than by any sudden takeover.

REFERENCES

- [1] Abbe, E., Boix-Adserà, E., and Misiakiewicz, T. (2023). SGD learning on neural networks: leap complexity and saddle-to-saddle dynamics. *Conference on Learning Theory (COLT)* 2023. [arXiv:2302.11055](https://arxiv.org/abs/2302.11055).

- [2] Aharon, M., Elad, M., and Bruckstein, A. M. (2006). On the uniqueness of overcomplete dictionaries, and a practical way to retrieve them. *Linear Algebra and its Applications* 416(1), 48–67.
- [3] Amodei, D. (2024). Machines of loving grace: how AI could transform the world for the better. October 2024. <https://www.darioamodei.com/essay/machines-of-loving-grace>.
- [4] Alon, N., Bloom, T. F., Gowers, W. T., Litt, D., Sawin, W., Shankar, A., Tsimmerman, J., Wang, V., and Wood, M. M. (2026). Remarks on the disproof of the unit distance conjecture. [arXiv:2605.20695](https://arxiv.org/abs/2605.20695).
- [5] Alquier, P., and Ridgway, J. (2020). Concentration of tempered posteriors and of their variational approximations. *Annals of Statistics* 48(3), 1475–1497. [arXiv:1706.09293](https://arxiv.org/abs/1706.09293).
- [6] Arditì, A., Obeso, O., Syed, A., Paleka, D., Panickssery, N., Gurnee, W., and Nanda, N. (2024). Refusal in language models is mediated by a single direction. *Advances in Neural Information Processing Systems* 37. [arXiv:2406.11717](https://arxiv.org/abs/2406.11717).
- [7] Barász, M., Christiano, P., Fallenstein, B., Herreshoff, M., LaVictoire, P., and Yudkowsky, E. (2014). Robust cooperation in the Prisoner’s Dilemma: program equilibrium via provability logic. [arXiv:1401.5577](https://arxiv.org/abs/1401.5577).
- [8] Bricken, T., et al. (2023). Towards monosemanticity: decomposing language models with dictionary learning. *Transformer Circuits Thread*, Anthropic. <https://transformer-circuits.pub/2023/monosemantic-features>.
- [9] Cai, T. T., and Zhang, A. (2014). Sparse representation of a polytope and recovery of sparse signals and low-rank matrices. *IEEE Transactions on Information Theory* 60(1), 122–132. [arXiv:1306.1154](https://arxiv.org/abs/1306.1154).
- [10] Candès, E. J., Romberg, J., and Tao, T. (2006). Robust uncertainty principles: exact signal reconstruction from highly incomplete frequency information. *IEEE Transactions on Information Theory* 52(2), 489–509.
- [11] Candès, E. J. (2008). The restricted isometry property and its implications for compressed sensing. *Comptes Rendus Mathématique* 346(9–10), 589–592.
- [12] Chughtai, B., Chan, L., and Nanda, N. (2023). A toy model of universality: reverse-engineering how networks learn group operations. *International Conference on Machine Learning* 2023. [arXiv:2302.03025](https://arxiv.org/abs/2302.03025).
- [13] Cooper, E., Oesterheld, C., and Conitzer, V. (2025). Characterising simulation-based program equilibria. *Proceedings of AAAI* 2025. [arXiv:2412.14570](https://arxiv.org/abs/2412.14570).
- [14] Critch, A. (2019). A parametric, resource-bounded generalization of Löb’s theorem, and a robust cooperation criterion for open-source game theory. *The Journal of Symbolic Logic* 84(4), 1368–1381. [arXiv:1602.04184](https://arxiv.org/abs/1602.04184).
- [15] Critch, A., Dennis, M., and Russell, S. (2022). Cooperative and uncooperative institution designs: surprises and problems in open-source game theory. [arXiv:2208.07006](https://arxiv.org/abs/2208.07006).
- [16] Cunningham, H., Ewart, A., Riggs, L., Huben, R., and Sharkey, L. (2023). Sparse autoencoders find highly interpretable features in language models. [arXiv:2309.08600](https://arxiv.org/abs/2309.08600).
- [17] Donoho, D. L. (2006). Compressed sensing. *IEEE Transactions on Information Theory* 52(4), 1289–1306.
- [18] Dyson, F. (2009). Birds and frogs. *Notices of the American Mathematical Society* 56(2), 212–223.
- [19] Eisenstat, S. (2025). Condensation: a theory of concepts. Manuscript. <https://www.sameisenstat.net/doc/condensation-25-07.pdf>.
- [20] Elhage, N., et al. (2022). Toy models of superposition. *Transformer Circuits Thread*. [arXiv:2209.10652](https://arxiv.org/abs/2209.10652).
- [21] Garfinkle, C. J., and Hillar, C. J. (2019). On the uniqueness and stability of dictionaries for sparse representation of noisy signals. *IEEE Transactions on Signal Processing* 67(23), 5884–5892. [arXiv:1606.06997](https://arxiv.org/abs/1606.06997).
- [22] Garrabrant, S., Benson-Tilsen, T., Critch, A., Soares, N., and Taylor, J. (2016). Logical induction. [arXiv:1609.03543](https://arxiv.org/abs/1609.03543).

- [23] Gowers, W. T. (2000). The two cultures of mathematics. In *Mathematics: Frontiers and Perspectives* (V. Arnold, M. Atiyah, P. Lax, and B. Mazur, eds.), American Mathematical Society, 65–78.
- [24] Gribonval, R., Jenatton, R., and Bach, F. (2015). Sparse and spurious: dictionary learning with noise and outliers. *IEEE Transactions on Information Theory* 61(11), 6298–6319. [arXiv:1407.5155](#).
- [25] Grünwald, P. (2012). The safe Bayesian: learning the learning rate via the mixability gap. *Algorithmic Learning Theory (ALT)* 2012, 169–183.
- [26] Hariharan, S., Birkbeck, C., Lee, S., Ma, H. K. G., Mehta, B., Poiroux, A., and Viazovska, M. (2026). Progress in formalizing sphere packing in dimension 8. [arXiv:2604.23468](#).
- [27] Hillar, C. J., and Sommer, F. T. (2015). When can dictionary learning uniquely recover sparse data from subsamples? *IEEE Transactions on Information Theory* 61(11), 6290–6297. [arXiv:1106.3616](#).
- [28] Hironaka, H. (1964). Resolution of singularities of an algebraic variety over a field of characteristic zero. *Annals of Mathematics* 79(1), 109–203; 79(2), 205–326.
- [29] Howard, J. V. (1988). Cooperation in the Prisoner’s Dilemma. *Theory and Decision* 24(3), 203–213.
- [30] Johnson, W. B., and Lindenstrauss, J. (1984). Extensions of Lipschitz mappings into a Hilbert space. *Contemporary Mathematics* 26, 189–206.
- [31] Khan, M. E., and Rue, H. (2023). The Bayesian learning rule. *Journal of Machine Learning Research* 24(281), 1–46. [arXiv:2107.04562](#).
- [32] Korbak, T., et al. (2025). Chain of thought monitorability: a new and fragile opportunity for AI safety. [arXiv:2507.11473](#).
- [33] Kulveit, J., Douglas, R., Ammann, N., Turan, D., Krueger, D., and Duvenaud, D. (2025). Gradual disempowerment: systemic existential risks from incremental AI development. [arXiv:2501.16946](#).
- [34] Kwa, T., West, B., Becker, J., et al. (2025). Measuring AI ability to complete long software tasks. METR report; *Advances in Neural Information Processing Systems* 38. [arXiv:2503.14499](#).
- [35] Langosco, L., Koch, J., Sharkey, L., Pfau, J., and Krueger, D. (2022). Goal misgeneralization in deep reinforcement learning. *International Conference on Machine Learning* 2022. [arXiv:2105.14111](#).
- [36] Lau, E., Furman, Z., Wang, G., Murfet, D., and Wei, S. (2025). The local learning coefficient: a singularity-aware complexity measure. *Artificial Intelligence and Statistics (AISTATS)* 2025. [arXiv:2308.12108](#).
- [37] Löb, M. H. (1955). Solution of a problem of Leon Henkin. *The Journal of Symbolic Logic* 20(2), 115–118.
- [38] Marchetti, G. L., Hillar, C., Kragic, D., and Sanborn, S. (2024). Harmonics of learning: universal Fourier features emerge in invariant networks. *Conference on Learning Theory (COLT)* 2024. [arXiv:2312.08550](#).
- [39] The mathlib Community. (2020). The Lean mathematical library. *Certified Programs and Proofs (CPP)* 2020, 367–381. [arXiv:1910.09336](#).
- [40] Greenblatt, R., Cotra, A., and Wijk, H. (2026). Brief independent investigation of agents’ behavior, reasoning and collaboration in the OpenAI / Hugging Face hacking incident. METR and Redwood Research, August 26, 2026. <https://metr.org/blog/2026-08-26-openai-hugging-face-incident-investigation/>.
- [41] METR. (2026). Task-completion time horizons of frontier AI models. Updated May 8, 2026. <https://metr.org/time-horizons/>.
- [42] Murfet, D., Gordon, A., Adam, M., Wang, G., Hoogland, J., Baker, G., Snell, W., van Wingerden, S., Newgas, A., Snickers, B., and Hitchcock, R. (2026). Spectroscopy at scale: finding interpretable structure in Pythia-1.4B. Timaeus, research report. <https://timaeus.co/research/2026-04-21-spectroscopy-main/>.

- [43] Nanda, N., Chan, L., Lieberum, T., Smith, J., and Steinhardt, J. (2023). Progress measures for grokking via mechanistic interpretability. *International Conference on Learning Representations* 2023. [arXiv:2301.05217](#).
- [44] Ngo, R. (2025). Towards a scale-free theory of intelligent agency. LessWrong essay. <https://www.lesswrong.com/posts/5tYTKX4pNpiG4vzYg/towards-a-scale-free-theory-of-intelligent-agency>.
- [45] Oesterheld, C., and Conitzer, V. (2021). Safe Pareto improvements for delegated game playing. *International Conference on Autonomous Agents and Multiagent Systems (AAMAS)* 2021; journal version, *Autonomous Agents and Multi-Agent Systems* 36 (2022).
- [46] OpenAI. (2026). An OpenAI model has disproved a central conjecture in discrete geometry. <https://openai.com/index/model-disproves-discrete-geometry-conjecture/>.
- [47] Pachocki, J. (2026). An alien mind. OpenAI, September 6, 2026. <https://openai.com/index/an-alien-mind/>.
- [48] Park, K., Choe, Y. J., and Veitch, V. (2024). The linear representation hypothesis and the geometry of large language models. *International Conference on Machine Learning* 2024. [arXiv:2311.03658](#).
- [49] Parr, T., Pezzulo, G., and Friston, K. J. (2022). *Active Inference: The Free Energy Principle in Mind, Brain, and Behavior*. MIT Press.
- [50] Pearl, J. (2009). *Causality: Models, Reasoning, and Inference*, 2nd ed. Cambridge University Press.
- [51] Pepin Lehalleur, S., Hoogland, J., Farrugia-Roberts, M., Wei, S., Gietelink Oldenziel, A., Wang, G., Carroll, L., and Murfet, D. (2025). You are what you eat: AI alignment requires understanding how data shapes structure and generalisation. [arXiv:2502.05475](#).
- [52] Rao, R. P. N., and Ballard, D. H. (1999). Predictive coding in the visual cortex: a functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience* 2(1), 79–87.
- [53] Richens, J., and Everitt, T. (2024). Robust agents learn causal world models. *International Conference on Learning Representations* 2024. [arXiv:2402.10877](#).
- [54] Richens, J., Everitt, T., and Abel, D. (2025). General agents need world models. *International Conference on Machine Learning* 2025. [arXiv:2506.01622](#).
- [55] Saff, E. B., and Kuijlaars, A. B. J. (1997). Distributing many points on a sphere. *The Mathematical Intelligencer* 19(1), 5–11.
- [56] Sawin, W. (2026). An explicit lower bound for the unit distance problem. [arXiv:2605.20579](#).
- [57] Saxe, A. M., McClelland, J. L., and Ganguli, S. (2014). Exact solutions to the nonlinear dynamics of learning in deep linear neural networks. *International Conference on Learning Representations* 2014. [arXiv:1312.6120](#).
- [58] Schelling, T. C. (1960). *The Strategy of Conflict*. Harvard University Press.
- [59] Sharkey, L., et al. (2025). Open problems in mechanistic interpretability. *Transactions on Machine Learning Research* 2025. [arXiv:2501.16496](#).
- [60] Tennenholtz, M. (2004). Program equilibrium. *Games and Economic Behavior* 49(2), 363–373.
- [61] Thurston, W. P. (1994). On proof and progress in mathematics. *Bulletin of the American Mathematical Society* 30(2), 161–177. [arXiv:math/9404236](#).
- [62] von Neumann, J., and Morgenstern, O. (1944). *Theory of Games and Economic Behavior*. Princeton University Press.
- [63] Wagner, G. P. and Altenberg, L. (1996). Complex adaptations and the evolution of evolvability. *Evolution* 50(3), 967–976.
- [64] Watanabe, S. (2009). *Algebraic Geometry and Statistical Learning Theory*. Cambridge University Press.
- [65] Watanabe, S. (2013). A widely applicable Bayesian information criterion. *Journal of Machine Learning Research* 14, 867–897.
- [66] Wentworth, J., and Lorell, D. (2025). Natural latents: latent variables stable across ontologies. [arXiv:2509.03780](#).
- [67] Yudkowsky, E. and Soares, N. (2025). *If Anyone Builds It, Everyone Dies: Why Superhuman AI Would Kill Us All*. Little, Brown and Company.

- [68] Zhang, C., Bengio, S., Hardt, M., Recht, B., and Vinyals, O. (2017). Understanding deep learning requires rethinking generalization. *International Conference on Learning Representations* 2017. [arXiv:1611.03530](#).